Hw 1

September 7, 2022

[1]:
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

1 Problem 1

Origin is qualitative. The rest are qualitative. (Note: I do not consider name to be a predictor)

[2]:
```
df = pd.read_csv("auto.csv")
```

1.1 Range and Mean and Standard Deviation for All data points

[3]:
```
df.iloc[:,:7].describe()
```

[3]:

	mpg	cylinders	displacement	weight	acceleration	year
count	397.000000	397.000000	397.000000	397.000000	397.000000	397.000000
mean	23.515869	5.458438	193.532746	2970.261965	15.555668	75.994962
std	7.825804	1.701577	104.379583	847.904119	2.749995	3.690005
min	9.000000	3.000000	68.000000	1613.000000	8.000000	70.000000
25%	17.500000	4.000000	104.000000	2223.000000	13.800000	73.000000
50%	23.000000	4.000000	146.000000	2800.000000	15.500000	76.000000
75%	29.000000	8.000000	262.000000	3609.000000	17.100000	79.000000
max	46.600000	8.000000	455.000000	5140.000000	24.800000	82.000000

1.2 Range Mean and Standard Dev for Datset excluding the 10th through 85th entries

[4]:
```
df.drop(df.index[10:85]).iloc[:,:7].describe()
```

[4]:
```
              mpg   cylinders  displacement       weight  acceleration  \
count  322.000000  322.000000    322.000000   322.000000    322.000000
mean    24.409317    5.378882    187.680124  2936.807453     15.700621
std      7.913357    1.657398    100.120925   810.987533      2.706436
min     11.000000    3.000000     68.000000  1649.000000      8.500000
25%     18.000000    4.000000    100.250000  2216.000000     14.000000
50%     23.900000    4.000000    145.500000  2797.500000     15.500000
75%     30.650000    6.000000    250.000000  3516.000000     17.275000
max     46.600000    8.000000    455.000000  4997.000000     24.800000

             year
count  322.000000
mean    77.130435
std      3.131849
min     70.000000
25%     75.000000
50%     77.000000
75%     80.000000
max     82.000000
```

1.3 Graphical Investigation of Predictors

[5]:
```
sns.pairplot(df)
```

[5]:
```
<seaborn.axisgrid.PairGrid at 0x1ba0bc8ea00>
```

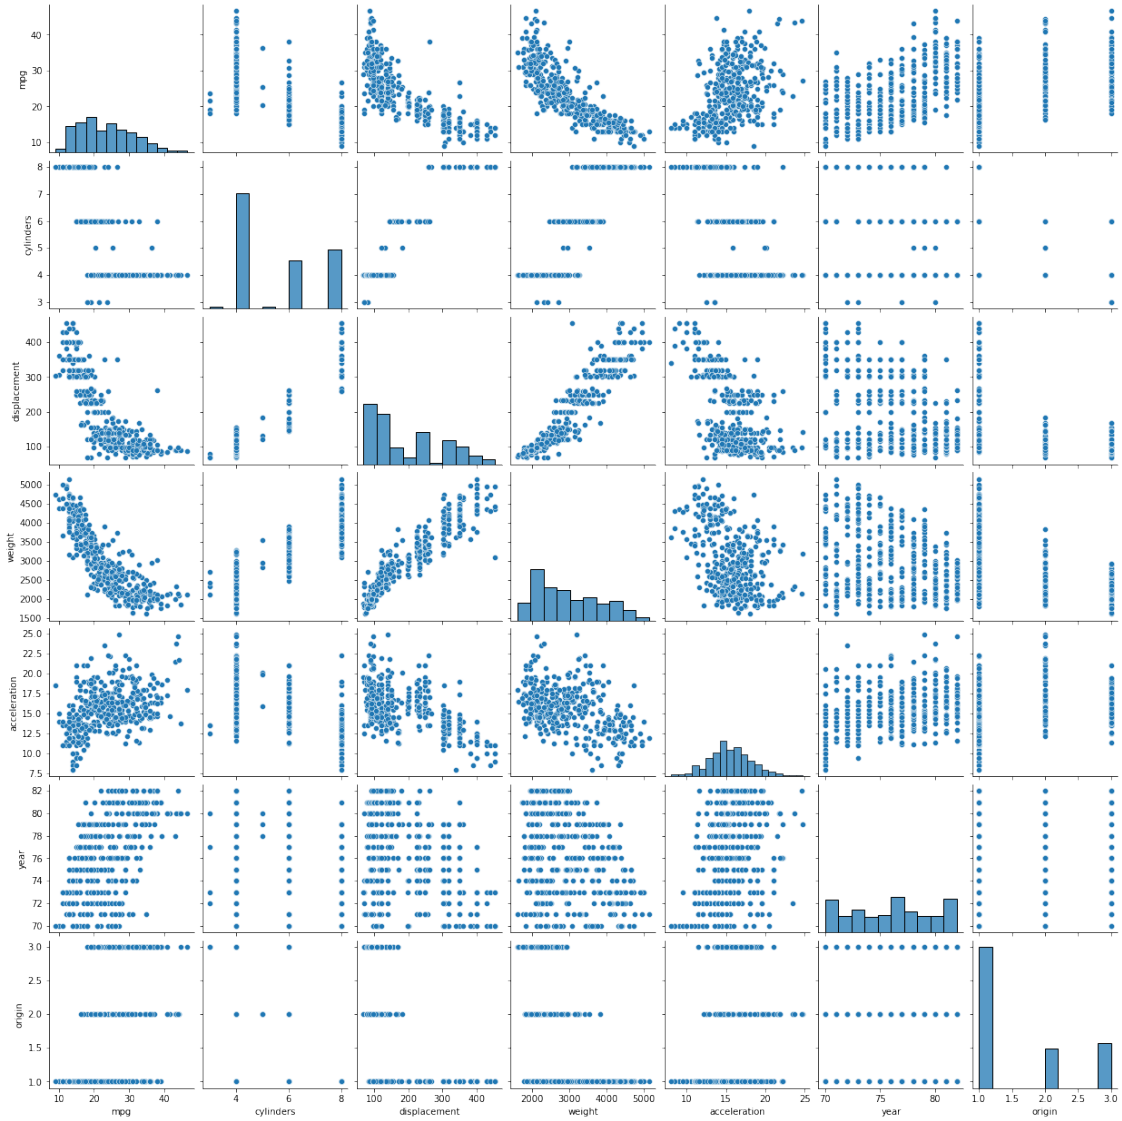

Some of the predictors are correlated with each other like displacement and weight. This could introduce confounding variable problem.

1.4 Correlation Matrix of Dataset

[6]:
```
plt.figure(figsize = (10,10))
sns.heatmap(df.corr(),annot = True)
```

[6]: <AxesSubplot:>

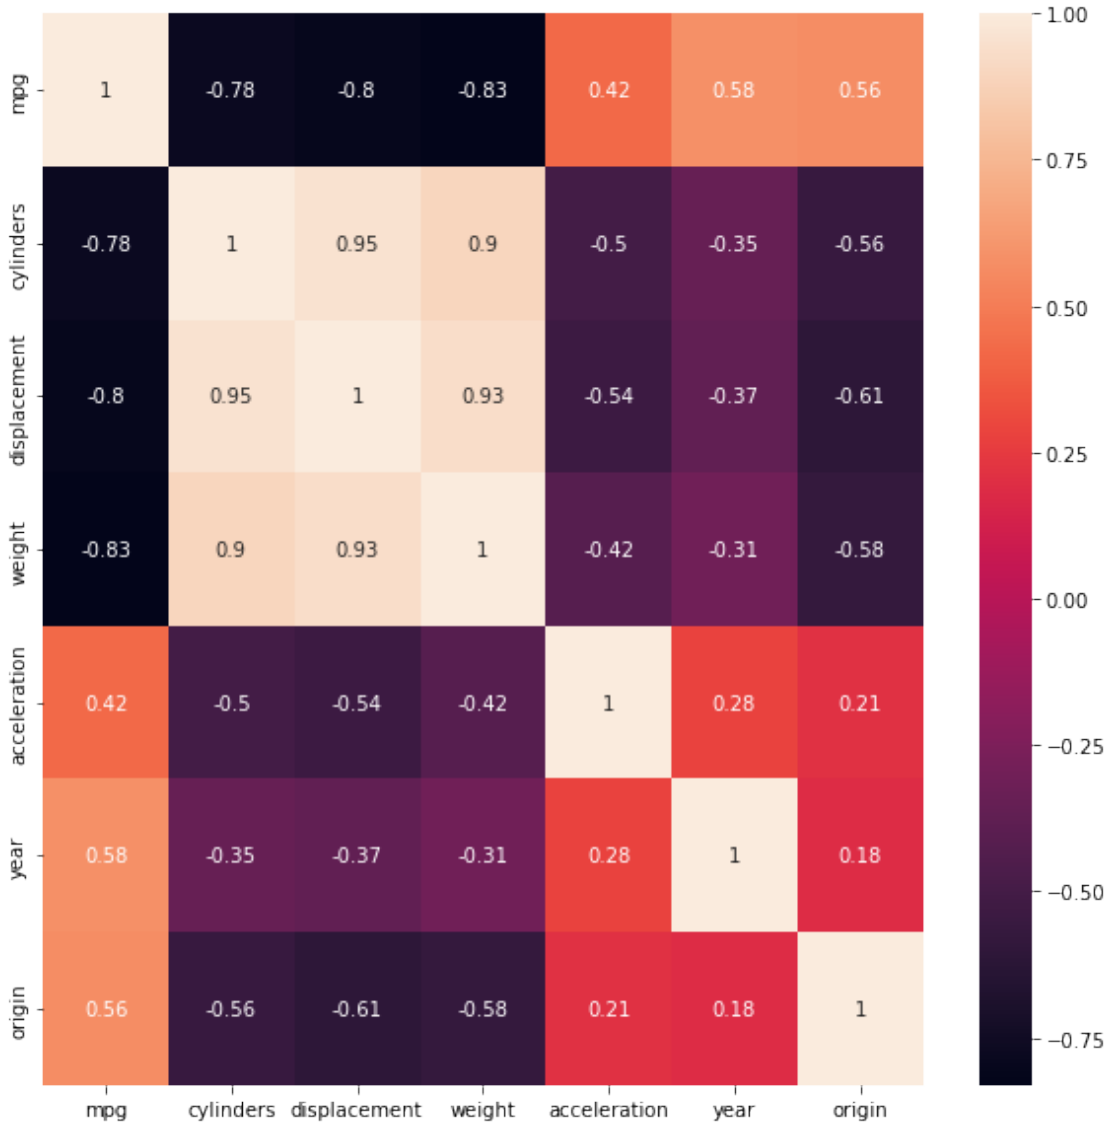

All Predictors seem to have good Correlation with MPG, though origin needs to be one hot encoded to be used.

2 Problem 2

[7]:
```
df = pd.read_csv('CodParasite.txt', sep="\t", index_col = 0)
df["log_intensity"] = np.log(1+ df["Intensity"])
```

[8]:
```
df
```

[8]:
```
         Intensity  Prevalence  Year  Depth  Weight  Length  Sex  Stage  Age  \
Sample
```

```
1             0.0           0  1999    220   148.0    26.0    0       0    0
2             0.0           0  1999    220   144.0    26.0    0       0    0
3             0.0           0  1999    220   146.0    27.0    0       0    0
4             0.0           0  1999    220   138.0    26.0    0       0    0
5             0.0           0  1999    220    40.0    17.0    0       0    0
…             …         …   …    …      …       …   …   …    …   …
1250         90.0           1  2001    228   224.0    31.0    1       1    2
1251        104.0           1  2001    140   690.0    43.0    2       1    3
1252        125.0           1  2001    140   754.0    44.0    2       1    3
1253        128.0           1  2001    140  1270.0    55.0    2       4    7
1254        257.0           1  2001    228   370.0    35.0    2       1    3

        Area  log_intensity
Sample
1          2       0.000000
2          2       0.000000
3          2       0.000000
4          2       0.000000
5          2       0.000000
…       …             …
1250       4       4.510860
1251       4       4.653960
1252       4       4.836282
1253       4       4.859812
1254       4       5.552960

[1254 rows x 11 columns]
```

[9]:
```
df.boxplot(column = ["log_intensity"], by = "Area", figsize = (10,10))
```

[9]:
```
<AxesSubplot:title={'center':'log_intensity'}, xlabel='Area'>
```

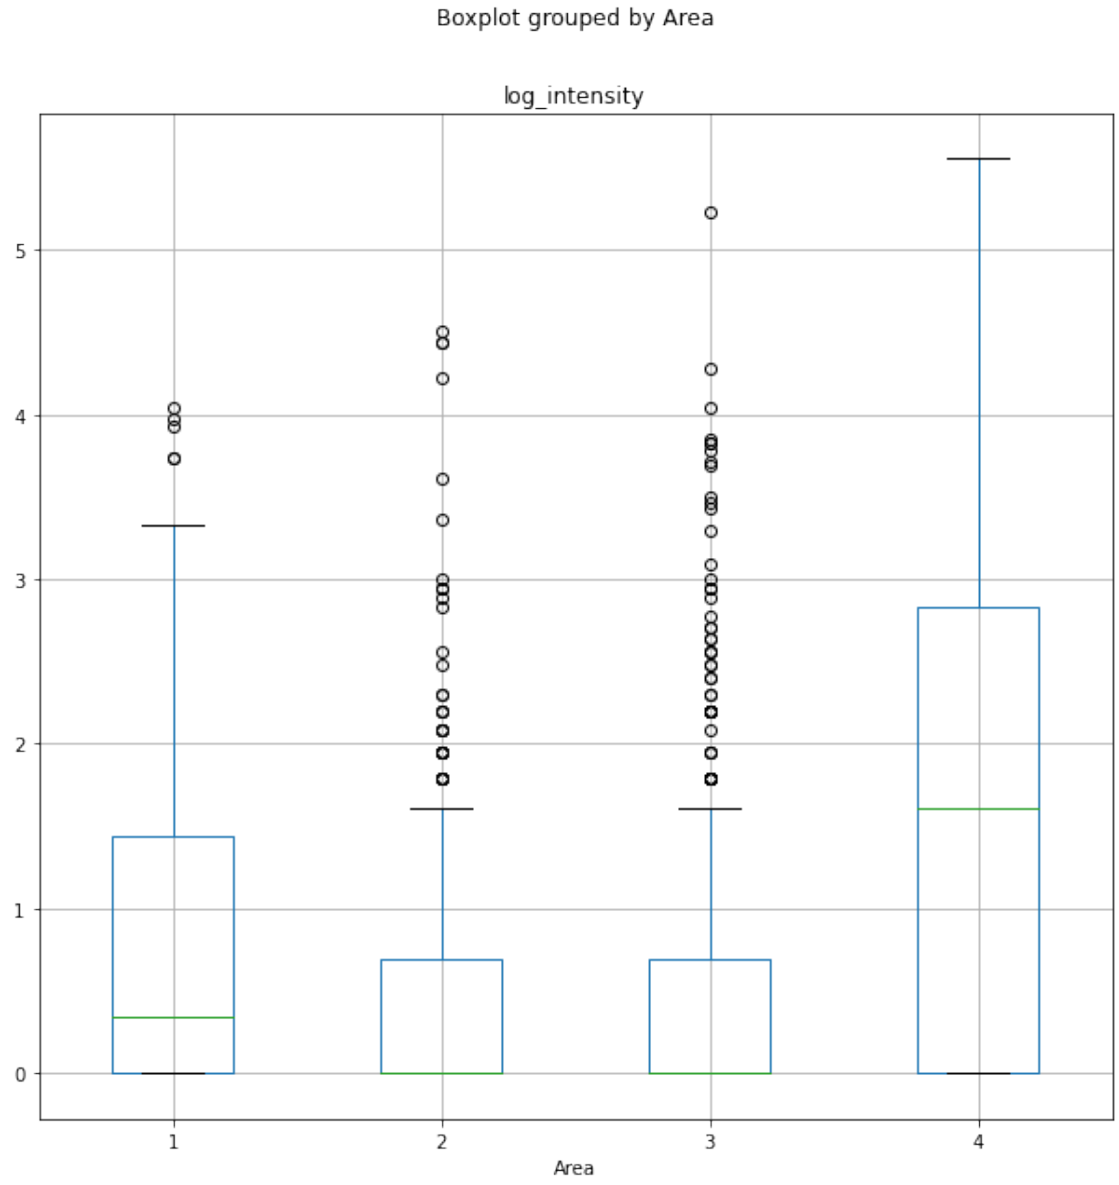

[10]:
```
df.boxplot(column = ["log_intensity"], by = "Sex", figsize = (10,10))
```

[10]: <AxesSubplot:title={'center':'log_intensity'}, xlabel='Sex'>

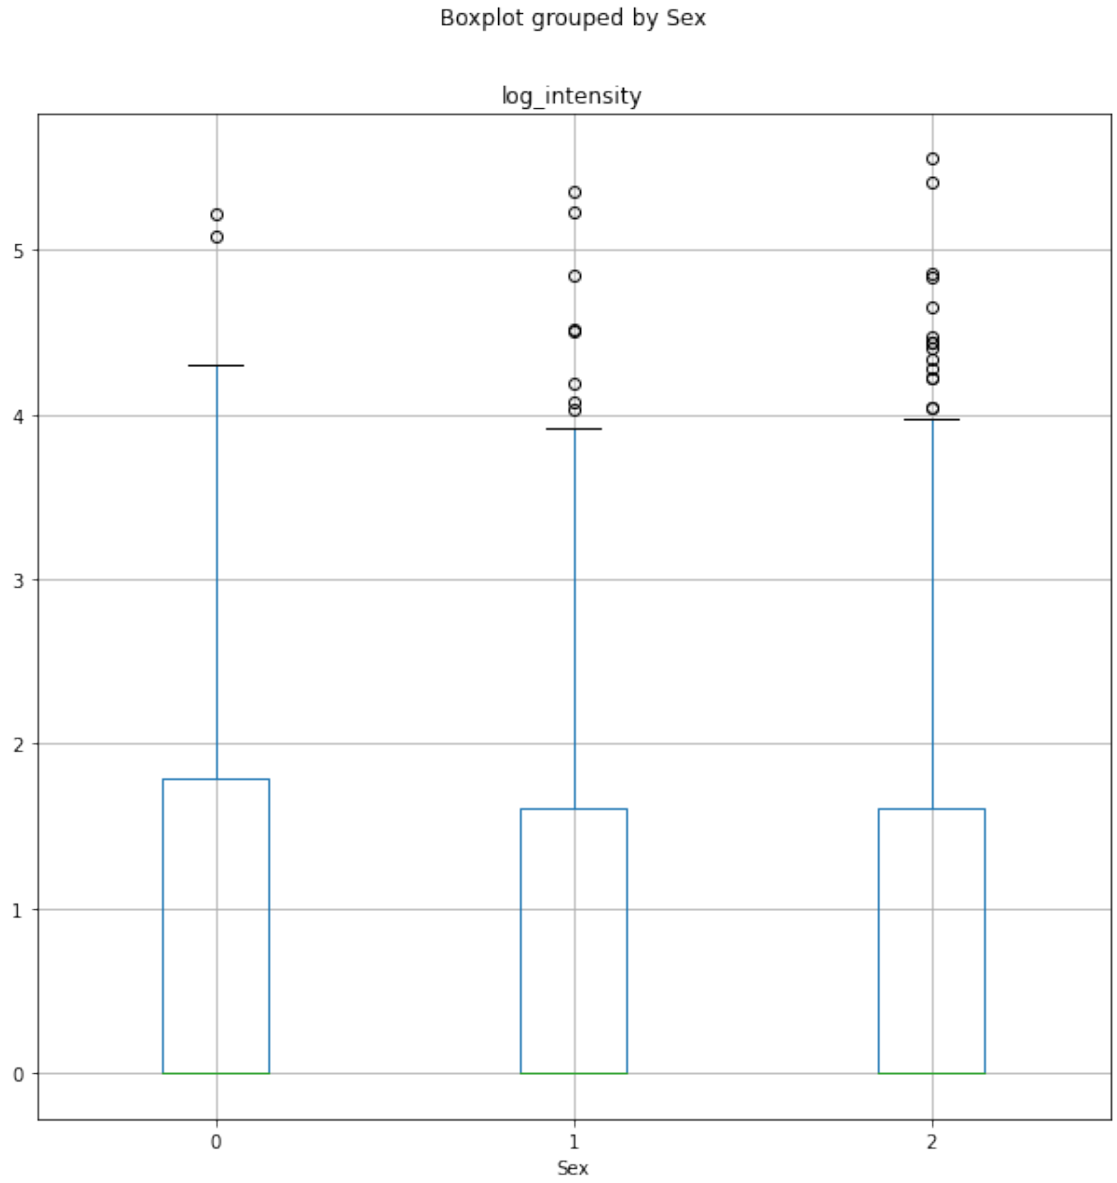

[11]:
```
df.boxplot(column = ["log_intensity"], by = "Stage", figsize = (10,10))
```

[11]: <AxesSubplot:title={'center':'log_intensity'}, xlabel='Stage'>

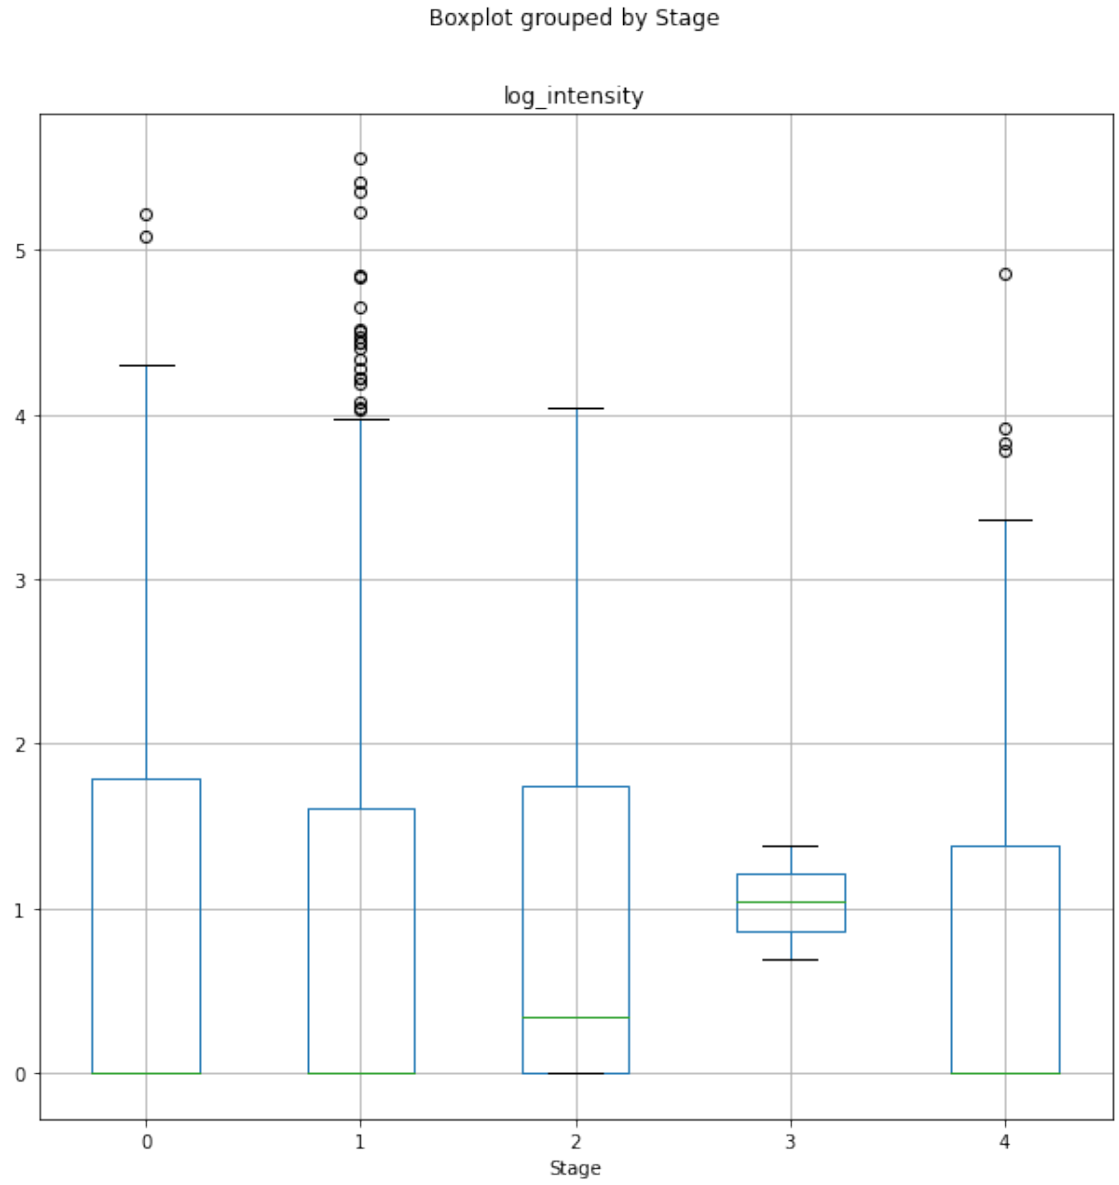

[12]:
```
df.boxplot(column = ["log_intensity"], by = "Year", figsize = (10,10))
```

[12]: <AxesSubplot:title={'center':'log_intensity'}, xlabel='Year'>

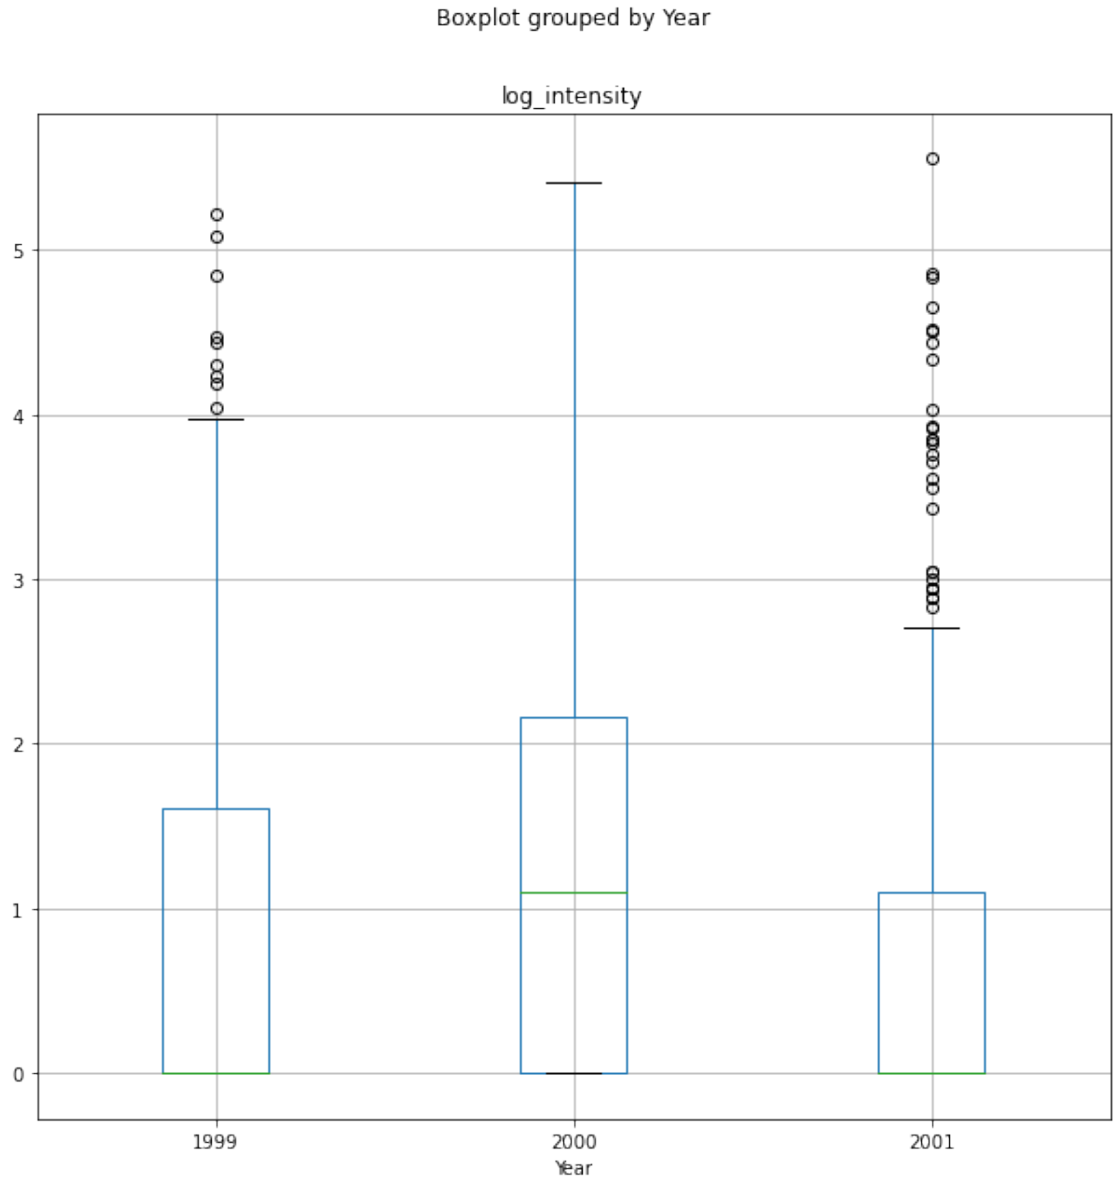

Area seems the most likely to be a good predictor of parasites since the distribution of each area is different from each other. All other box plots look similar to each other within the class.

3 Problem 3

[13]:
```
df = pd.read_csv('Owls.txt', sep="\t")
df
```

[13]:

	Nest	FoodTreatment	SexParent	ArrivalTime	SiblingNegotiation	\
0	AutavauxTV	Deprived	Male	22.25	4	
1	AutavauxTV	Satiated	Male	22.38	0	

```
2    AutavauxTV      Deprived      Male        22.53                  2
3    AutavauxTV      Deprived      Male        22.56                  2
4    AutavauxTV      Deprived      Male        22.61                  2
..          …           …       …          …                  …
594    Yvonnand      Deprived    Female        27.25                  7
595    Yvonnand      Deprived      Male        28.45                  5
596    Yvonnand      Deprived    Female        28.86                 15
597    Yvonnand      Deprived      Male        29.21                 10
598    Yvonnand      Satiated    Female        29.23                  0

     BroodSize  NegPerChick
0            5     0.800000
1            5     0.000000
2            5     0.400000
3            5     0.400000
4            5     0.400000
..         …          …
594          7     1.000000
595          7     0.714286
596          7     2.142857
597          7     1.428571
598          7     0.000000

[599 rows x 7 columns]
```

[14]:
```
df.boxplot(column = ["NegPerChick"], by = "SexParent", figsize = (10,10))
```

[14]:
```
<AxesSubplot:title={'center':'NegPerChick'}, xlabel='SexParent'>
```

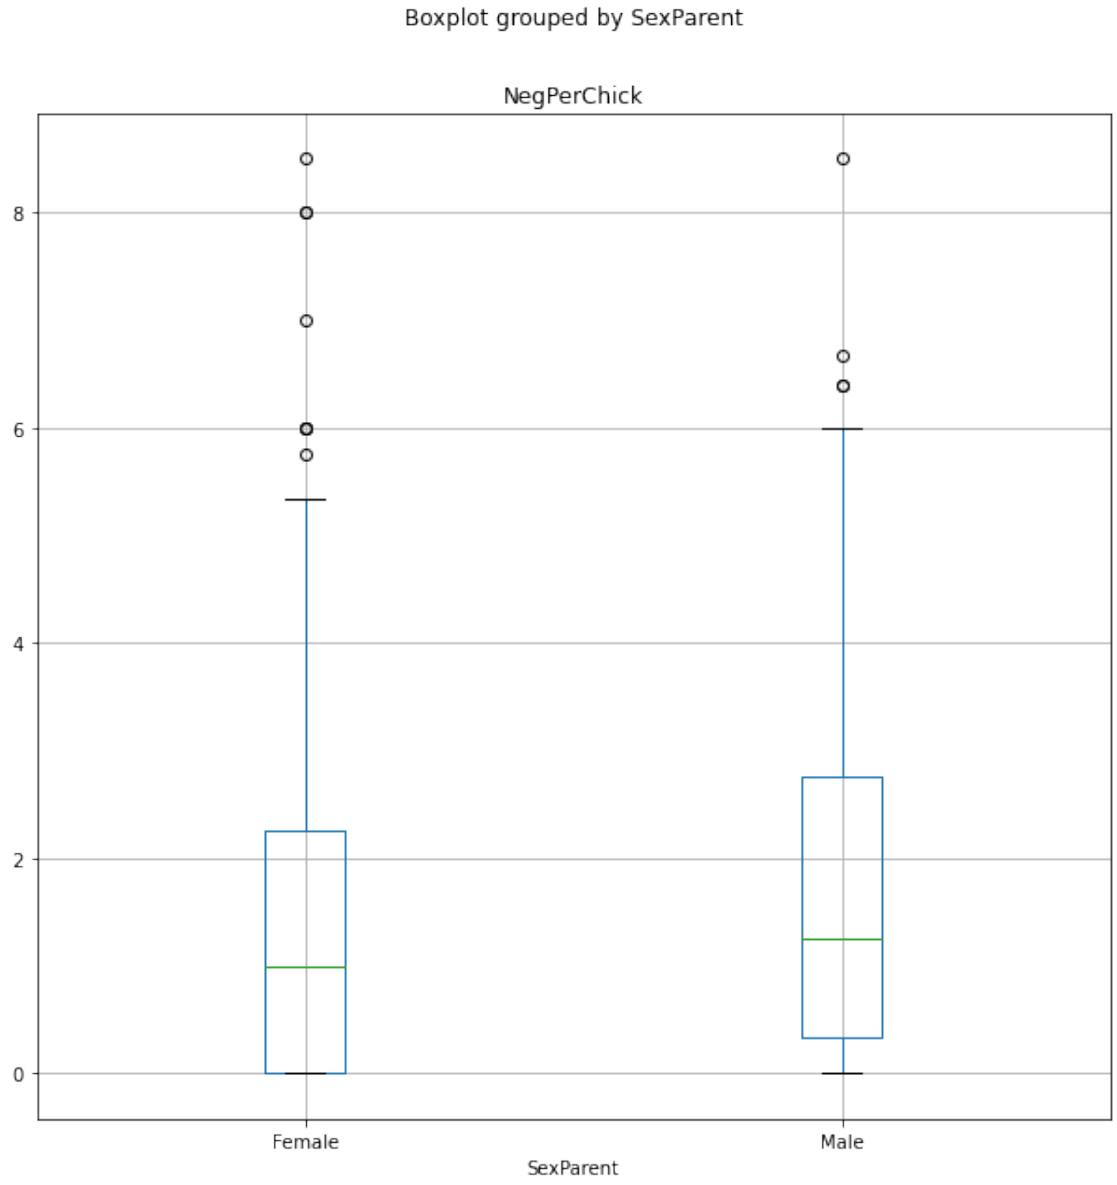

[15]:
```
df.boxplot(column = ["NegPerChick"], by = "FoodTreatment", figsize = (10,10))
```

[15]: <AxesSubplot:title={'center':'NegPerChick'}, xlabel='FoodTreatment'>

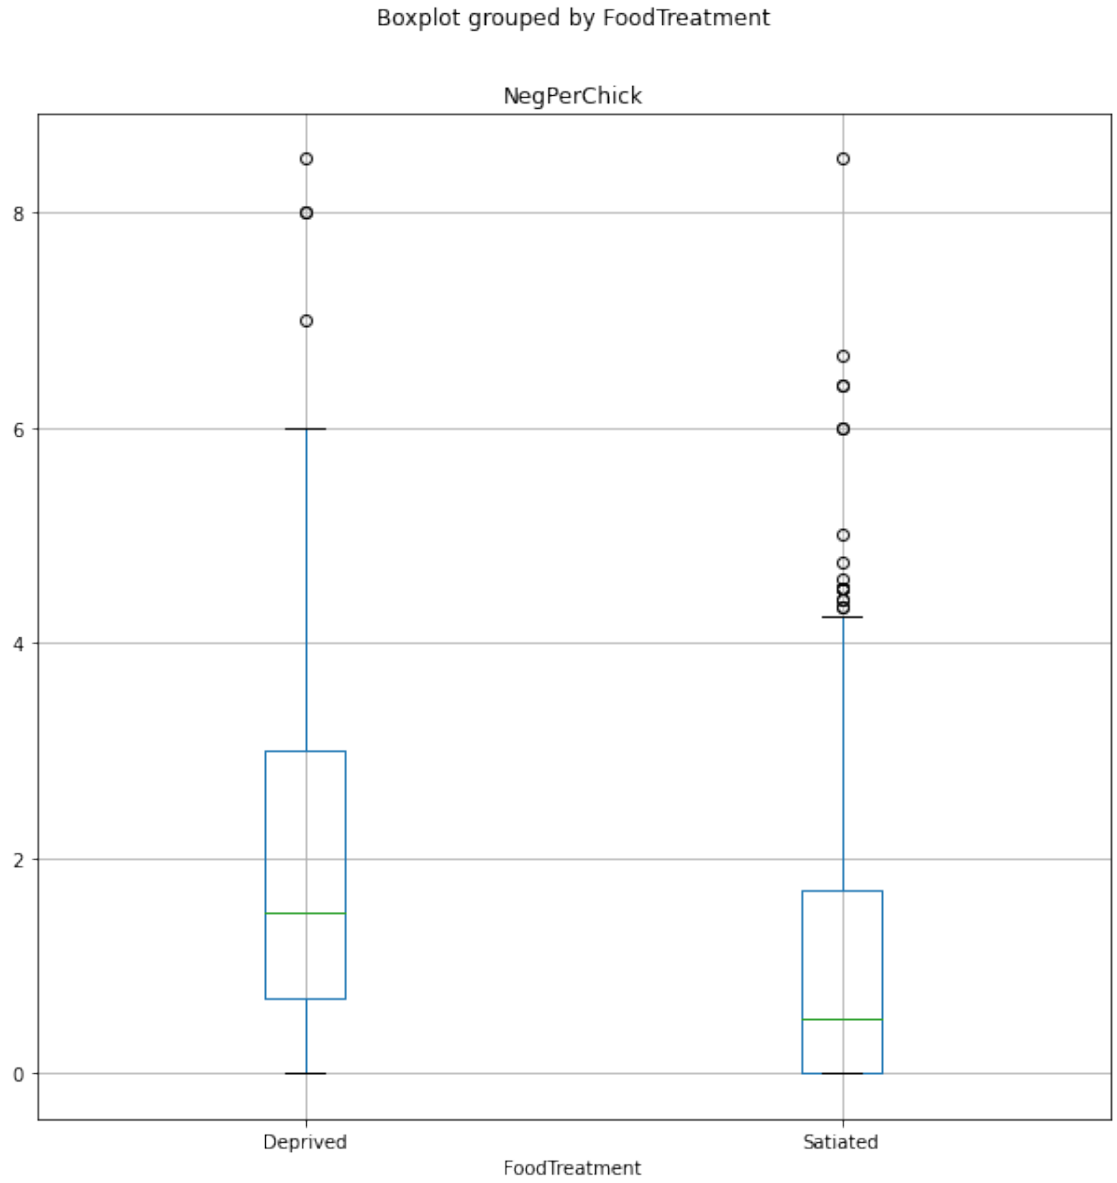

[16]:
```
df.boxplot(column = ["NegPerChick"], by = ["SexParent","FoodTreatment"],␣
↪figsize = (10,10))
```

[16]:
```
<AxesSubplot:title={'center':'NegPerChick'}, xlabel='[SexParent,
FoodTreatment]'>
```

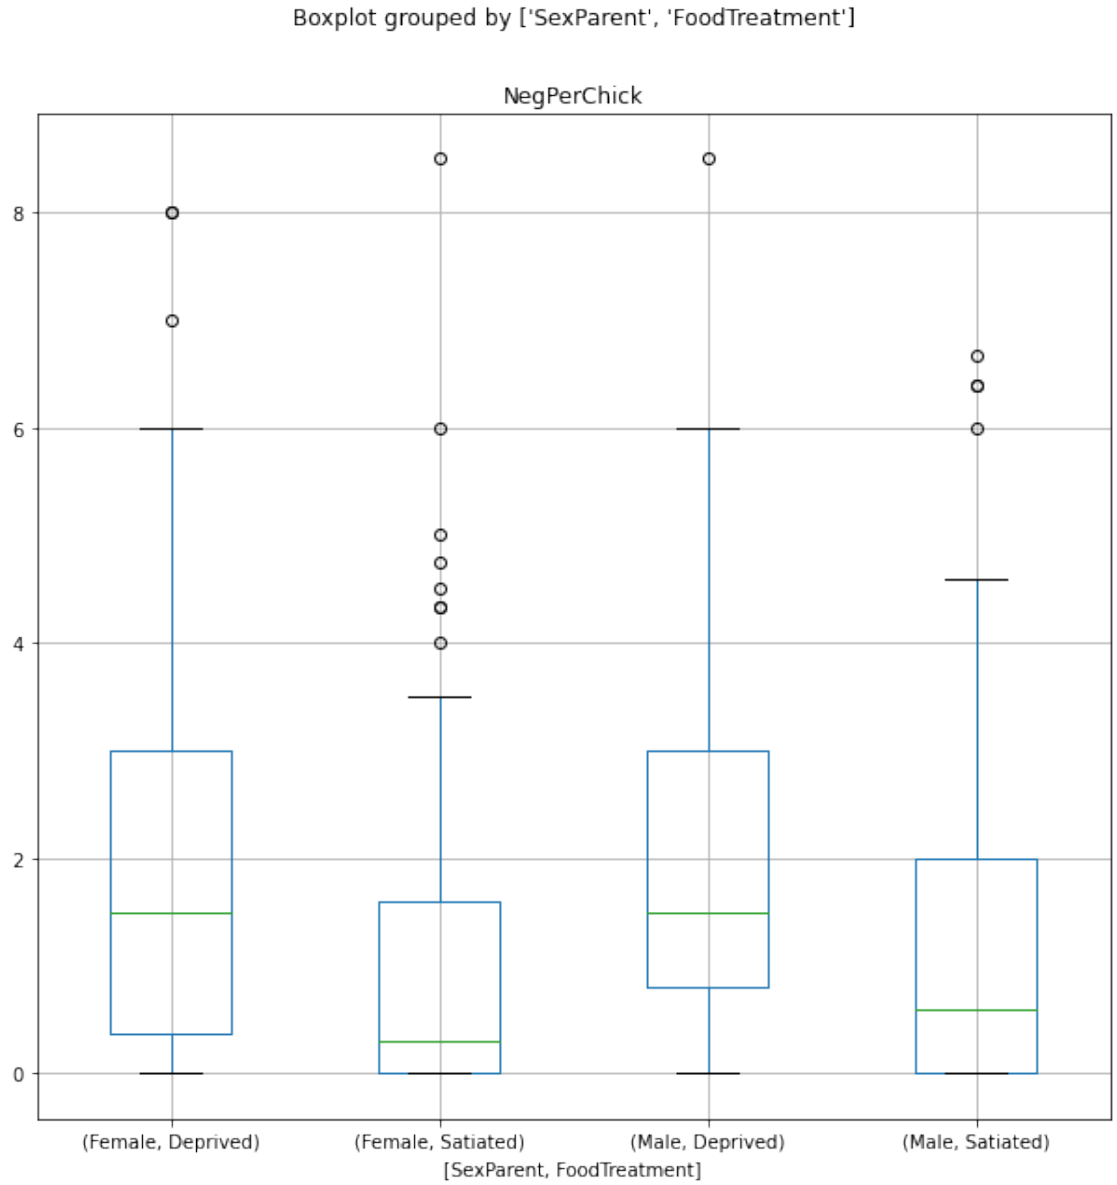

[17]:
```
df.boxplot(column = ["NegPerChick"], by = "Nest", figsize = (30,10))
```

[17]:
```
<AxesSubplot:title={'center':'NegPerChick'}, xlabel='Nest'>
```

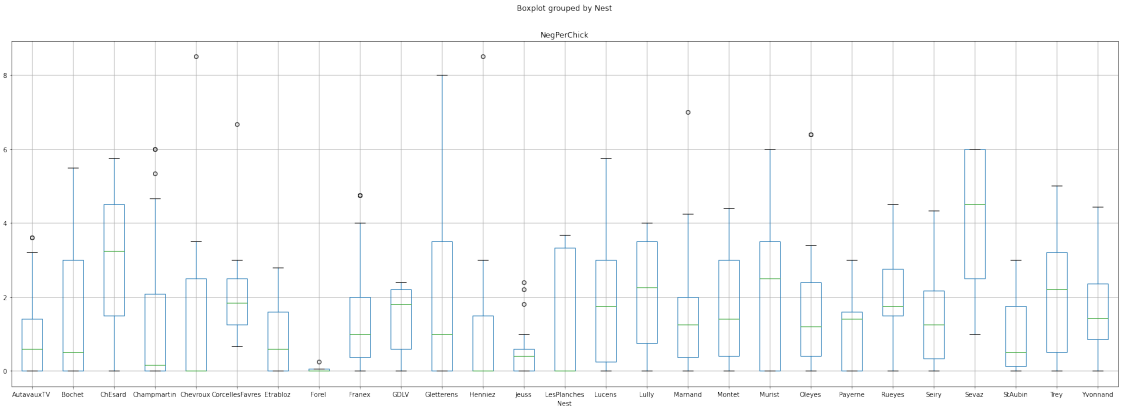

There do seem to be differences across different nesting sites as the distributions differ. This could also be due to the data having less points per nesting site thus increasing the variance.

[18]:
```
df.plot.scatter(x="NegPerChick", y = "SiblingNegotiation", figsize = (10,10))
```

[18]:
```
<AxesSubplot:xlabel='NegPerChick', ylabel='SiblingNegotiation'>
```

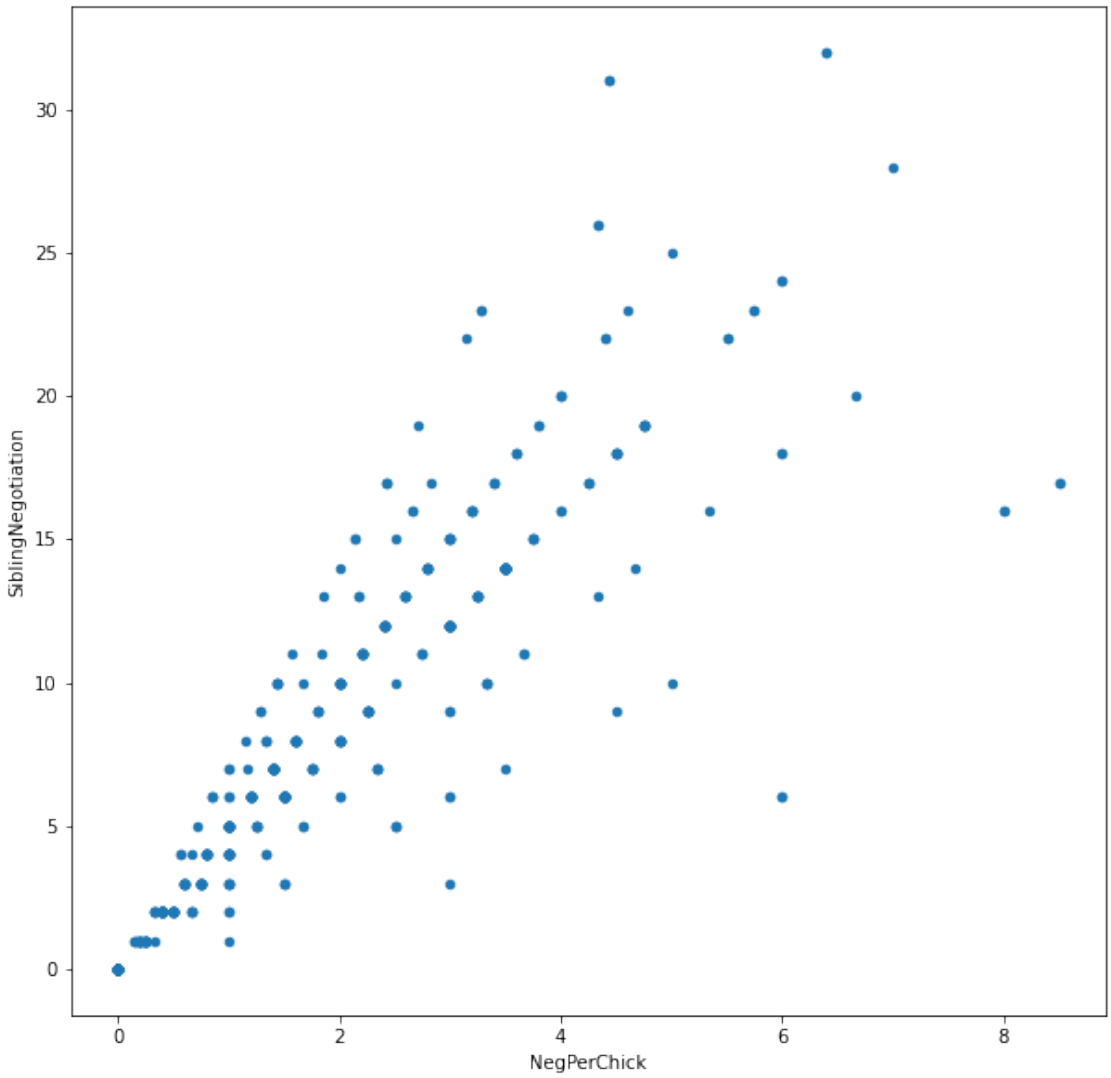

There seems to be some correlation with increasing variance as NegPerChick increases

[19]:
```
df["log_sibling"] = np.log(1+df["SiblingNegotiation"])
```

[20]:
```
df.plot.scatter(x="log_sibling", y = "ArrivalTime", figsize = (10,10))
```

[20]:
```
<AxesSubplot:xlabel='log_sibling', ylabel='ArrivalTime'>
```

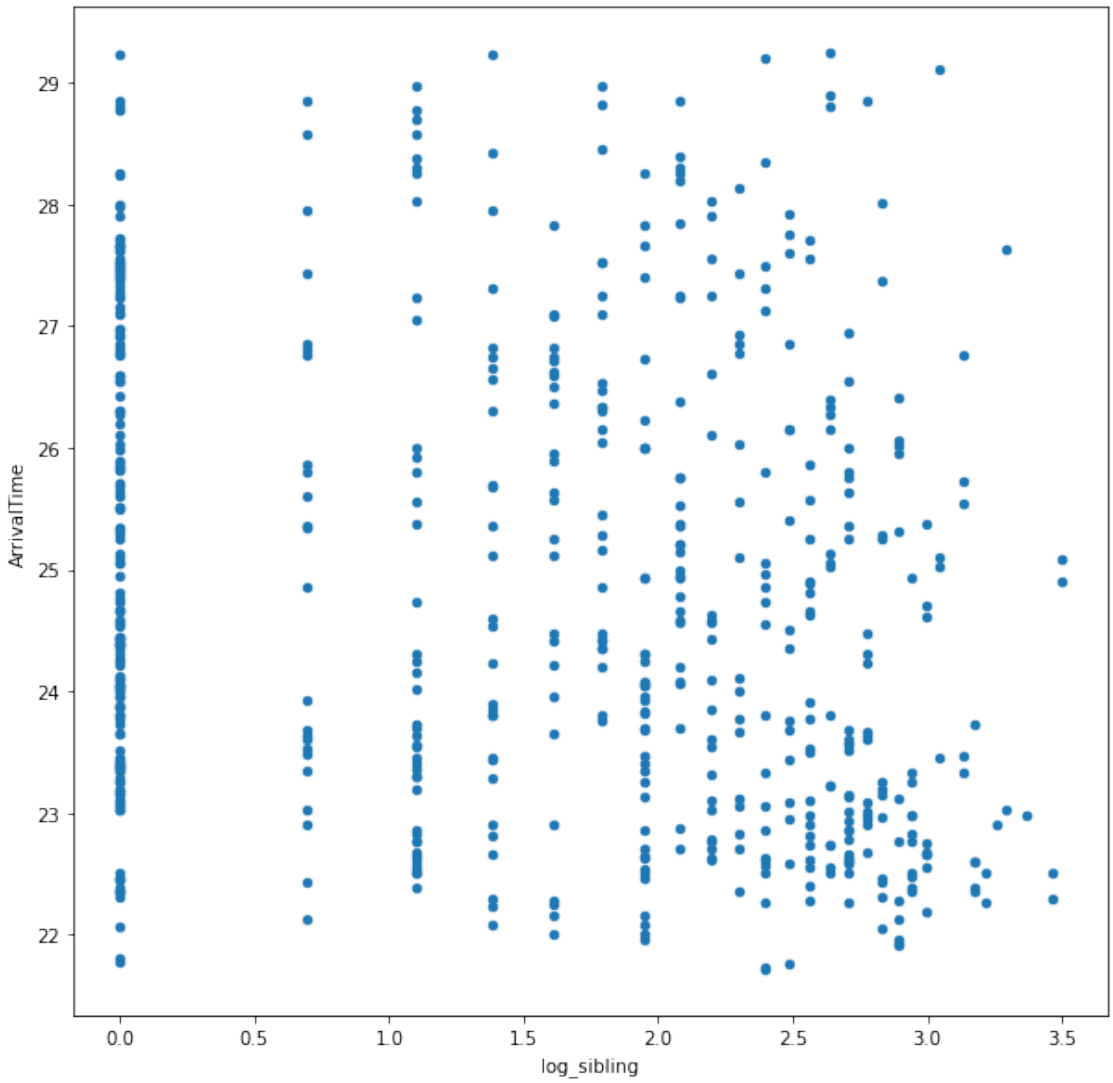
ArrivalTime
log_sibling
29
28
27
26
25
24
23
22
0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5