

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import PowerTransformer
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegressionCV
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [2]: df = pd.read_csv("magic04.csv")
df = df.drop_duplicates()
df = df.rename(columns = {"class": "Class"})
X = df.iloc[:,10]
y = df.iloc[:,10]
Class = y.Class.unique()
le = LabelEncoder()
y = pd.DataFrame(le.fit_transform(np.ravel(y)), columns = y.columns)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state = 101)
scaler = StandardScaler()
scaler.fit(X_train)
X_train = pd.DataFrame(scaler.transform(X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
pt = PowerTransformer(method = "yeo-johnson")
pt.fit(X_train)
X_train = pd.DataFrame(pt.transform(X_train), columns = X.columns)
X_test = pd.DataFrame(pt.transform(X_test), columns = X.columns)
```

Logistic Regression

```
In [3]: Logit = LogisticRegressionCV(cv=10, scoring='accuracy', n_jobs=-1, max_iter = 100)
Logit.fit(X_train, np.ravel(y_train))
y_pred = Logit.predict(X_test)
```

```
In [4]: pd.DataFrame(data = Logit.coef_, columns = X.columns)
```

```
Out[4]:
```

	fLength	fWidth	fSize	fConc	fConc1	fAsym	fM3Long	fM3Trans	fAlpha	fDist
0	1.383869	-0.189867	1.455469	1.421105	0.218006	-0.049338	-0.292599	-0.018409	1.586697	0.037657

Columns with Higher Coefficients are More Important (fLength, fSize, fConc, fAlpha). Others with low coefficient are less important.

```
In [5]: print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred, target_names=Class))
```

```
[[2292 201]
 [ 421 867]]
```

	precision	recall	f1-score	support
g	0.84	0.92	0.88	2493
h	0.81	0.67	0.74	1288
accuracy			0.84	3781
macro avg	0.83	0.80	0.81	3781
weighted avg	0.83	0.84	0.83	3781

```
In [6]: from sklearn import metrics
```

```
In [7]: fpr, tpr, _ = metrics.roc_curve(y_test, Logit.predict_proba(X_test)[:,1])
```

```
In [8]: metrics.auc(fpr, tpr)
```

```
Out[8]: 0.8745026446721628
```


```
In [9]: from sklearn.metrics import accuracy_score
```

```
In [10]: accuracy_score(y_test, y_pred)
```

```
Out[10]: 0.8354932557524465
```

```
In [11]: auc_logit = pd.DataFrame(data = {"fpr":fpr,"tpr":tpr})
```

```
In [12]: plt.figure(figsize = (15,10))
plot = sns.lineplot(x='fpr', y='tpr', data=auc_logit).set(title='AUC ROC LOGIT')
```



LDA

```
In [13]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
In [14]: LDA = LinearDiscriminantAnalysis()
```

```
In [15]: LDA.fit(X_train, np.ravel(y_train))
y_pred = LDA.predict(X_test)
```

```
In [16]: print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred, target_names=Class))
```

```
[[2290 203]
 [ 439 849]]
```

	precision	recall	f1-score	support
g	0.84	0.92	0.88	2493
h	0.81	0.66	0.73	1288
accuracy			0.83	3781
macro avg	0.82	0.79	0.80	3781
weighted avg	0.83	0.83	0.83	3781

```
In [17]: fpr, tpr, _ = metrics.roc_curve(y_test, LDA.predict_proba(X_test)[:,1])
```

```
In [18]: metrics.auc(fpr, tpr)
```


```
Out[18]: 0.8740526268583089
```

```
In [19]: accuracy_score(y_test, y_pred)
```

```
Out[19]: 0.8302036498280878
```

```
In [20]: auc_LDA = pd.DataFrame(data = {"fpr":fpr,"tpr":tpr})
```

```
In [21]: plt.figure(figsize = (15,10))
plot = sns.lineplot(x='fpr', y='tpr', data=auc_LDA).set(title='AUC ROC LDA')
```



KNN

```
In [22]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [ ]: from sklearn.model_selection import GridSearchCV
from sklearn.metrics import make_scorer
```

```
scoring = ("AUC", "roc_auc", "Accuracy": make_scorer(accuracy_score))

# Setting refit='AUC', refits an estimator on the whole dataset with the
# parameter setting that has the best cross-validated AUC score.
# That estimator is made available at 'gs.best_estimator_' along with
# parameters like 'gs.best_score_', 'gs.best_params_' and
# 'gs.best_index_'.
gs = GridSearchCV(
    KNeighborsClassifier(),
    param_grid={"n_neighbors": range(1, 21), "weights": ["uniform", "distance"], "p": [1, 2]},
    scoring=scoring,
    refit="AUC",
    return_train_score=True,
    n_jobs = -1,
    cv = 10,
    verbose = 3
)
gs.fit(X_train, np.ravel(y_train))
results = gs.cv_results_
```

```
In [ ]: gs.best_params_
```

```
In [23]: KNN = KNeighborsClassifier(n_neighbors=20, p = 2, weights = "distance")
KNN.fit(X_train, np.ravel(y_train))
y_pred = KNN.predict(X_test)
```

```
In [24]: print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred, target_names=Class))
```

```
[[2421  72]
 [ 457 831]]
```

	precision	recall	f1-score	support
g	0.84	0.97	0.90	2493
h	0.92	0.65	0.76	1288
accuracy			0.86	3781
macro avg	0.88	0.81	0.83	3781
weighted avg	0.87	0.86	0.85	3781

```
In [25]: fpr, tpr, _ = metrics.roc_curve(y_test, KNN.predict_proba(X_test)[:,1])
```

```
In [26]: metrics.auc(fpr, tpr)
```

```
Out[26]: 0.9197065447850254
```

```
In [27]: accuracy_score(y_test, y_pred)
```

```
Out[27]: 0.8600899233007141
```

```
In [28]: auc_KNN = pd.DataFrame(data = {"fpr":fpr,"tpr":tpr})
```

```
In [29]: plt.figure(figsize = (15,10))
plot = sns.lineplot(x='fpr', y='tpr', data=auc_KNN).set(title='AUC ROC KNN')
```



Linear SVM

```
In [30]: from sklearn.svm import SVC
```

```
In [ ]: C_range = np.logspace(-2, 10, 13)
param_grid = dict(C=C_range)
```

```
In [ ]: from sklearn.model_selection import StratifiedShuffleSplit
```

```
In [ ]: cv = StratifiedShuffleSplit(n_splits=30, test_size=0.2, random_state=101)
```

```
In [ ]: grid = GridSearchCV(SVC(kernel = "linear", max_iter=10000, probability=True), param_grid=param_grid, cv=cv, n_jobs=-1)
```

```
In [ ]: grid.fit(X_train, np.ravel(y_train))
```

```
In [ ]: results = grid.cv_results_
```

```
In [ ]: grid.best_params_
```

```
In [31]: LSVC = SVC(kernel = "linear", max_iter=10000, C = 0.1, probability = True)
LSVC.fit(X_train, np.ravel(y_train))
y_pred = LSVC.predict(X_test)
```

C:\Users\matth\anaconda3\lib\site-packages\sklearn\svm_base.py:255: ConvergenceWarning: Solver terminated early (max_iter=10000). Consider pre-processing your data with StandardScaler or MinMaxScaler.

warnings.warn('Solver terminated early (max_iter=%i).')

```
In [32]: print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred, target_names=Class))
```

```
[[2329 164]
 [ 450 838]]
```

	precision	recall	f1-score	support
g	0.84	0.93	0.88	2493
h	0.84	0.65	0.73	1288
accuracy			0.84	3781
macro avg	0.84	0.79	0.81	3781
weighted avg	0.84	0.84	0.83	3781

```
In [33]: fpr, tpr, _ = metrics.roc_curve(y_test, LSVC.predict_proba(X_test)[:,1])
```

```
In [34]: metrics.auc(fpr, tpr)
```

```
Out[34]: 0.8726321900077983
```

```
In [35]: accuracy_score(y_test, y_pred)
```

```
Out[35]: 0.8376090981221899
```

```
In [36]: auc_LSVC = pd.DataFrame(data = {"fpr":fpr,"tpr":tpr})
```

```
In [37]: plt.figure(figsize = (15,10))
plot = sns.lineplot(x='fpr', y='tpr', data=auc_LSVC).set(title='AUC ROC LINEAR SVC')
```



Radial SVC

```
In [ ]: gamma_range = np.logspace(-9, 3, 13)
param_grid = dict(gamma=gamma_range, C=C_range)
```

```
In [ ]: cv = StratifiedShuffleSplit(n_splits=3, test_size=0.2, random_state=101)
grid = GridSearchCV(SVC(probability = True), param_grid=param_grid, cv=cv, n_jobs = -1, verbose = 3)
```

```
In [ ]: grid.fit(X_train, np.ravel(y_train))
```

```
In [ ]: results = grid.cv_results_
```

```
In [ ]: grid.best_params_
```

```
In [38]: RSVC = SVC(C = 100, gamma = 0.1, probability = True)
RSVC.fit(X_train, np.ravel(y_train))
y_pred = RSVC.predict(X_test)
```

```
In [39]: print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred, target_names=Class))
```

```
[[2389 104]
 [ 318 970]]
```

	precision	recall	f1-score	support
g	0.88	0.96	0.92	2493
h	0.90	0.75	0.82	1288
accuracy			0.89	3781
macro avg	0.89	0.86	0.87	3781
weighted avg	0.89	0.89	0.89	3781

```
In [40]: fpr, tpr, _ = metrics.roc_curve(y_test, RSVC.predict_proba(X_test)[:,1])
```

```
In [41]: metrics.auc(fpr, tpr)
```

```
Out[41]: 0.9290217578162955
```

```
In [42]: accuracy_score(y_test, y_pred)
```

```
Out[42]: 0.8883893149960328
```

```
In [43]: auc_RSVC = pd.DataFrame(data = {"fpr":fpr,"tpr":tpr})
```

```
In [44]: plt.figure(figsize = (15,10))
plot = sns.lineplot(x='fpr', y='tpr', data=auc_RSVC).set(title='AUC ROC RADIAL SVC')
```

